

• Horizontal Scaling

- Load Balancing required
- **RESILIENT**
- Network calls (RPC)
- Data inconsistency
- **Scales well as users increase**


• Vertical Scaling

- N/A
- Single point of failure
- **Inter- process communications**
- **Consistent Data**
- Hardware limit

Consistent Hashing

Request Id $\rightarrow h(r1)$

$1 / M$



Server 0

$h1(0) \% M$

Server 1

$H1(1) \% M$

Server 2

$H1(2) \% M$

Server 3

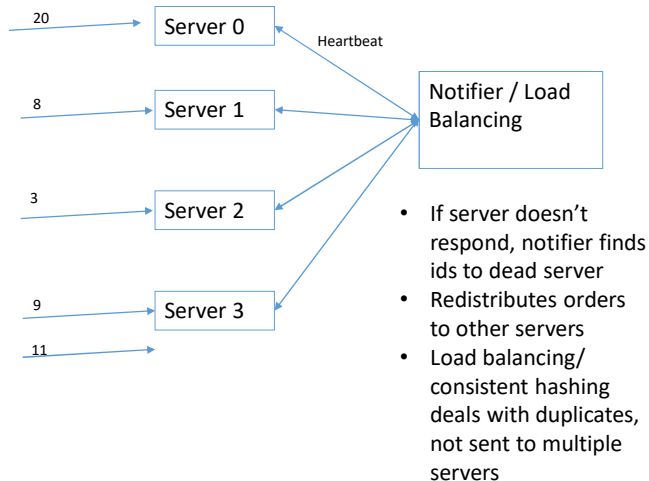
$H1(3) \% M$

Server M

$H1(M) \% M$

- Assign server to positions
- Hash request ID, and find closest server
- 2nd hash table, taking into account virtual servers, to load balance equally
- Each servers maps to multiple points
- Works well only when many servers

Message/ Task Queues



Features
 1) Assignment / notification
 2) Load balancing
 3) Heartbeat
 4) persistence

ID	Contents	Done
1	Ham	Y
2	Cheese	N
3	Plain	N
4

Monolithic vs MicroServices

• Monolithic

- Can have many machines
- Simpler to maintain
- Less Complex
- Don't need to duplicate for setting up tests, connections
- Procedure calls faster, not RPC
- Deployments are complicated, have to be monitored every time
- Single point of failure, have to restart everything instead of at few points

• MicroService

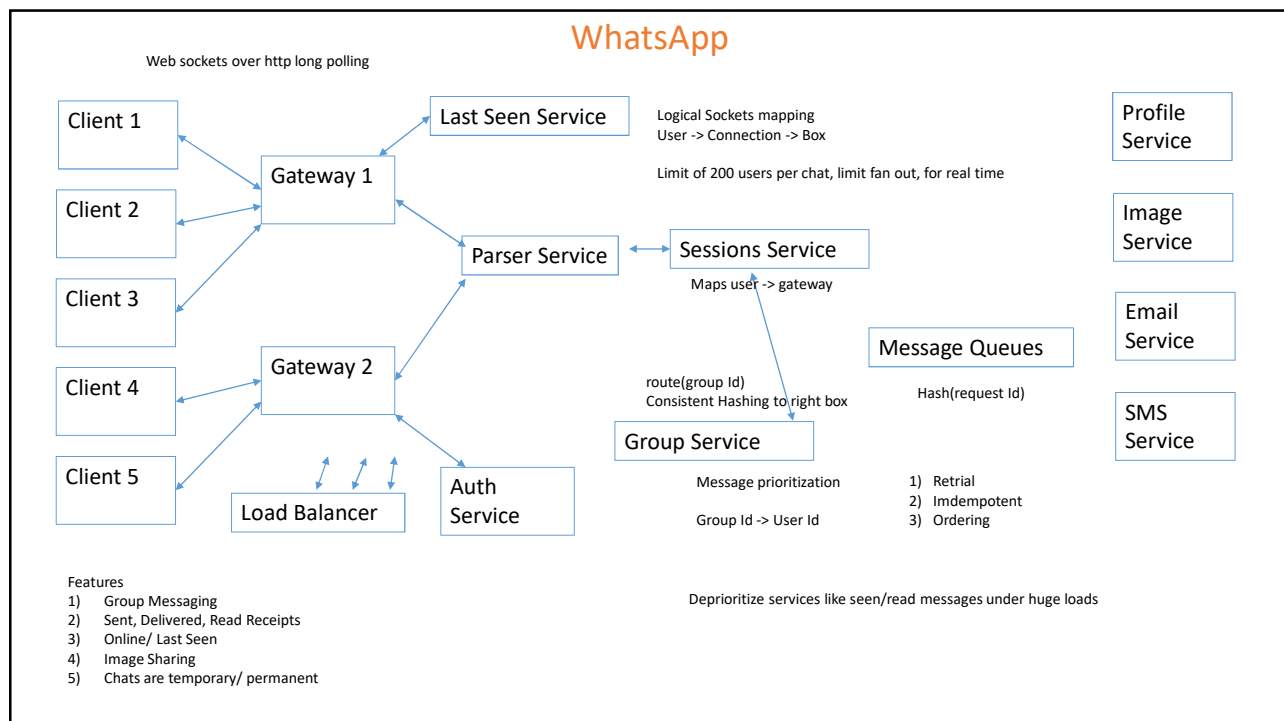
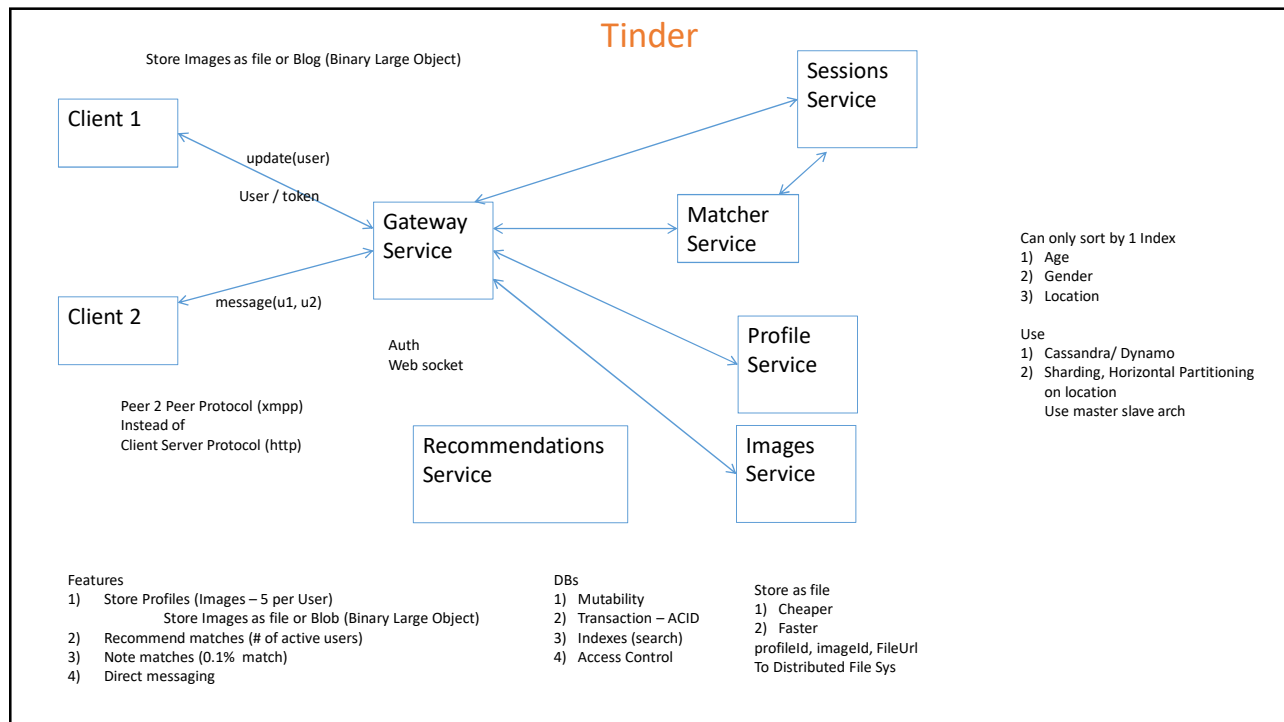
- Can have little machines
- Easier to scale
- Easier for new team members
- Parallel development is easier
- Fewer parts are hidden when deploying
- Not easy to design, needs smart architects

Database Sharding

- Consistency
- Availability
- Can shard by userId, location
- Problems
- Joins across shards, expensive
- Shards are inflexible in number, can use consistent hashing (Memcached)
- Create index on shards
- Master -> Slave Arch
- Writes to master, reads to slaves, slave becomes master if needed

Single point of failure

- More Nodes
- Master – Slave
- Regions
- Multiple Load Balancers (Gateway)
- Client -> DNS -> LBs -> Nodes -> DBs



Distributed Caching

Why

- 1) Avoid Network Call
 - 2) Avoid computations
 - 3) Reduce DB load
- If cache closer to server, faster, but can be inconsistent between servers (in memory)
 - Global cache (Redis), can recover when failing, and can scale independently
 - Write-through
 - Update cache first, then update DB
 - Can wait and send to DB in both, for noncritical data, for saving network calls
 - Not practical for multiple caches
 - Write-back (performance issues)
 - Update DB first, then make entry in cache/ or invalidate
 - When hit cache on GET, then go to db, and also update cache
 - Hit cache, invalidate entry if it is there
 - Hybrid
 - If not critical info, write to cache, wait, and take entries in bulk and write to DB
- 1) Eviction Policy
 - 2) Thrashing – constantly inputting and outputting to cache without using results
 - 3) Consistency

API Design

Items of importance

- 1) Naming
- 2) Parameters deficiency
 - 1) More params for efficiency only
- 3) Response object simplicity
- 4) Return specific errors messages, for expected errors

Design API request

- 1) www.webiste.com/chat_messaging/getAdmins
 - 1) POST
 - 2) Request Object
 - 3) Response
- 2) No side effects. If many flags, should break down into separate functions
 - 1) Doing everything
 - 2) atomicity
- 3) When response is huge
 - 1) Pagination - Break response into multiple responses, but not stateless
 - 2) Fragmentation
- 4) Do you want perfect consistency?

NoSQL

1) SQL

- 1) requires joins which are expensive, no way to efficiently normalize

2) NoSQL


- 1) Flexible schema
- 2) insertions and retrievals require the whole blob
- 3) Horizontal partitioning built in, build for scale
- 4) Build for aggregation, finding metrics

3) NoSQL – disadvantages



- 1) Not built for updates (delete and insert), Consistency, ACID not guaranteed
- 2) Not read optimized
- 3) Relations are not implicit
- 4) Joins are hard

4) Cassandra

- 1) Have n instances
- 2) Writes to m servers, where pos through $\text{pos} + m - 1$
- 3) Quorum factor of x. For reads, x servers have to return the same value




Distributed Consensus



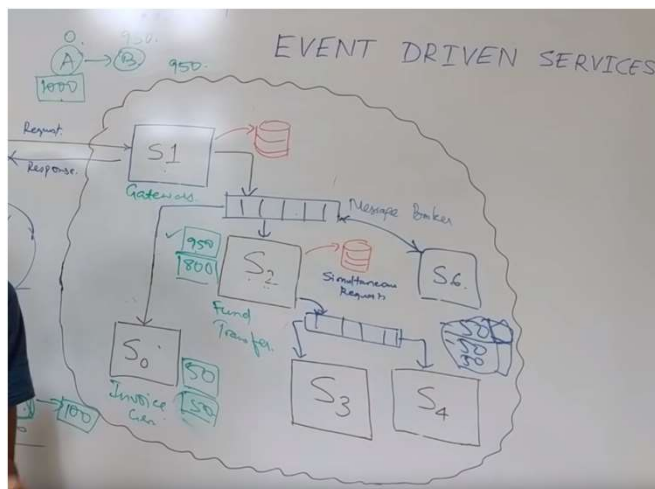
1) Bottleneck to Master DB

- 1) Add slave node, prevent writes to it, only reads
- 2) Or make both masters
- 2) Add 2 slaves
- 3) 2 Phase commit, 3 phase commit, MVCC, SAGA


Pub Sub



- 1) Use queue
- 2) Disadvantage – not enough consistency by default



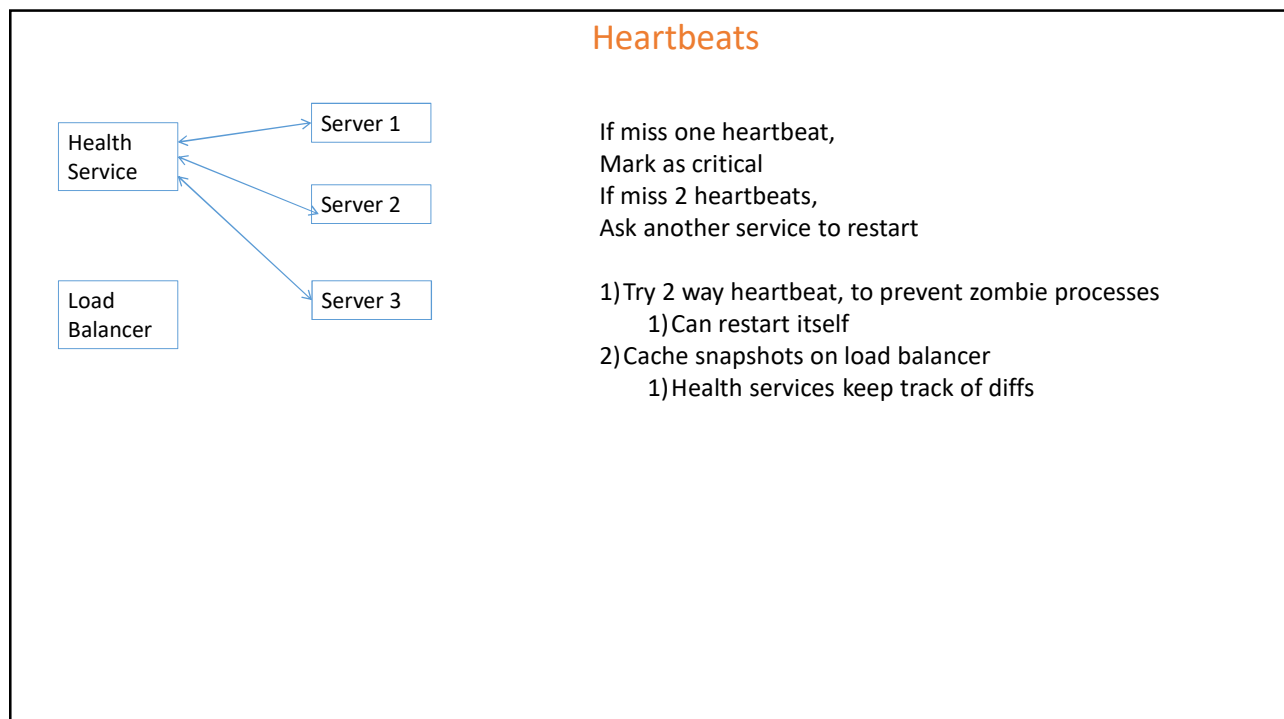
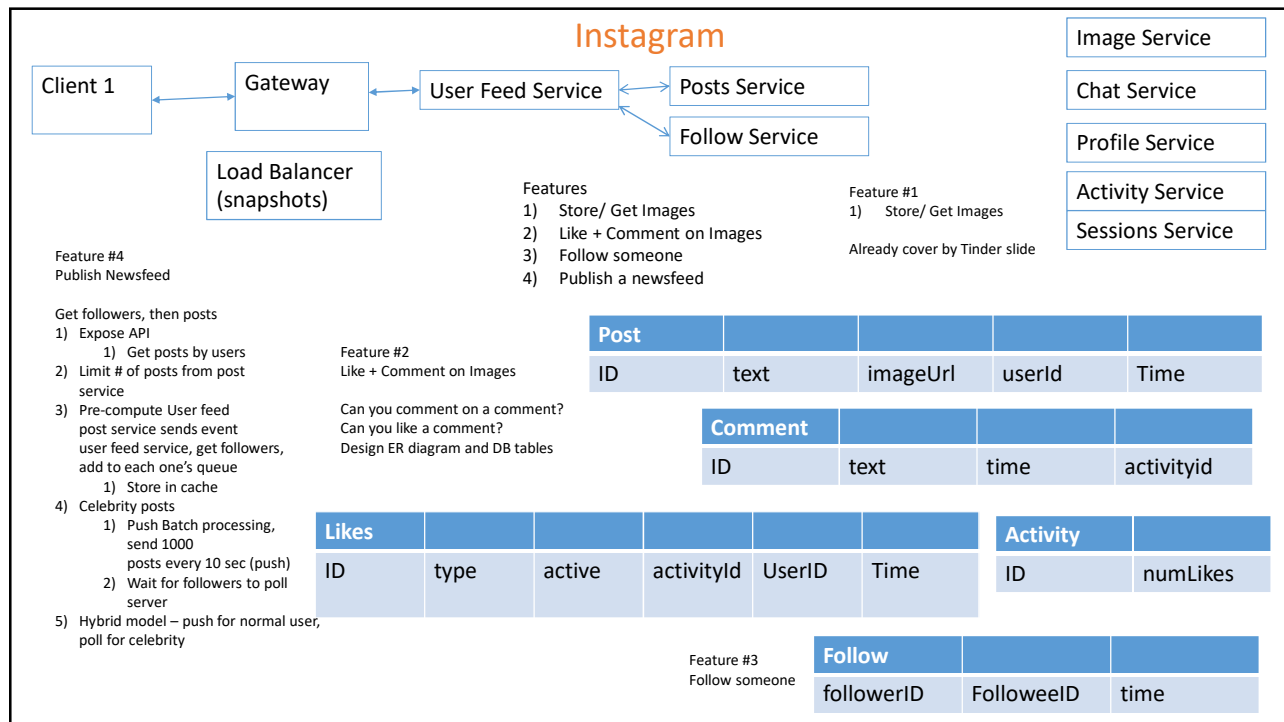
Event Driven Systems




- 1) Availability (but lower consistency)
- 2) Easy Roll-backs
- 3) Replacement

Problems

- 1) Consistency
- 2) N/A to Gateways
- 3) Less control of responses
- 4) Compaction
- 5) Hidden Slow
- 6) Fixing
 - 1) Replay from Start
 - 2) Diff based
 - 3) Undo



Tips



- 1) Don't go into detail prematurely
 - 1) Look for first point to go into detail
 - 2) Database -> ER diagrams
- 2) Do not have a set architecture in mind
- 3) KISS
 - 1) Keep It Simple Stupid
- 4) Form your thoughts
 - 1) Make points without justifications
 - 2) i.e. choosing a SQL database
- 5) Be Tech Aware