

- **Horizontal Scaling**

- Load Balancing required
- **RESILIENT**
- Network calls (RPC)
- Data inconsistency
- **Scales well as users increase**


- **Vertical Scaling**

- N/A
- Single point of failure
- **Inter- process communications**
- **Consistent Data**
- Hardware limit

Consistent Hashing

Request Id $\rightarrow h(r1)$

$1 / M$



Server 0

$h1(0) \% M$

Server 1

$H1(1) \% M$

Server 2

$H1(2) \% M$

Server 3

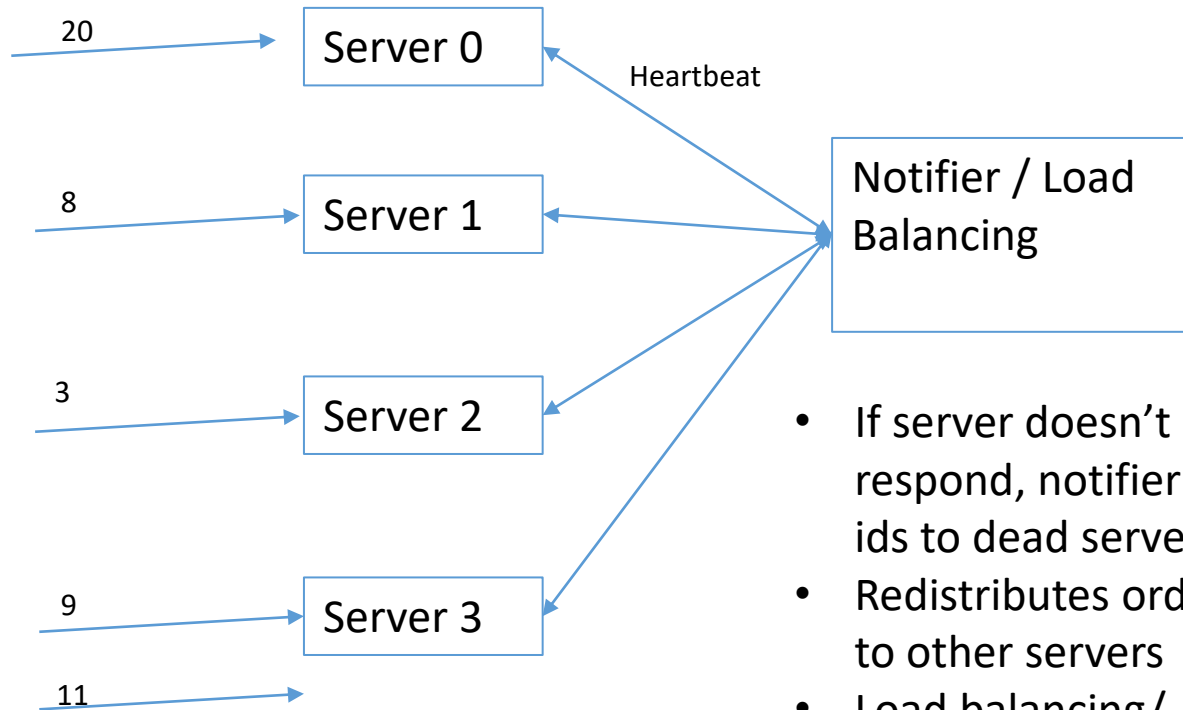
$H1(3) \% M$

Server M

$H1(M) \% M$

- Assign server to positions
- Hash request ID, and find closest server
- 2nd hash table, taking into account virtual servers, to load balance equally
- Each servers maps to multiple points
- Works well only when many servers

Message/ Task Queues



- If server doesn't respond, notifier finds ids to dead server
- Redistributes orders to other servers
- Load balancing/ consistent hashing deals with duplicates, not sent to multiple servers

ID	Contents	Done
1	Ham	Y
2	Cheese	N
3	Plain	N
4

Features

- 1) Assignment / notification
- 2) Load balancing
- 3) Heartbeat
- 4) persistence

Monolithic vs MicroServices

- **Monolithic**

- Can have many machines
- Simpler to maintain
- Less Complex
- Don't need to duplicate for setting up tests, connections
- Procedure calls faster, not RPC
- Deployments are complicated, have to be monitored every time
- Single point of failure, have to restart everything instead of at few points

- **MicroService**

- Can have little machines
- Easier to scale
- Easier for new team members
- Parallel development is easier
- Fewer parts are hidden when deploying
- Not easy to design, needs smart architects

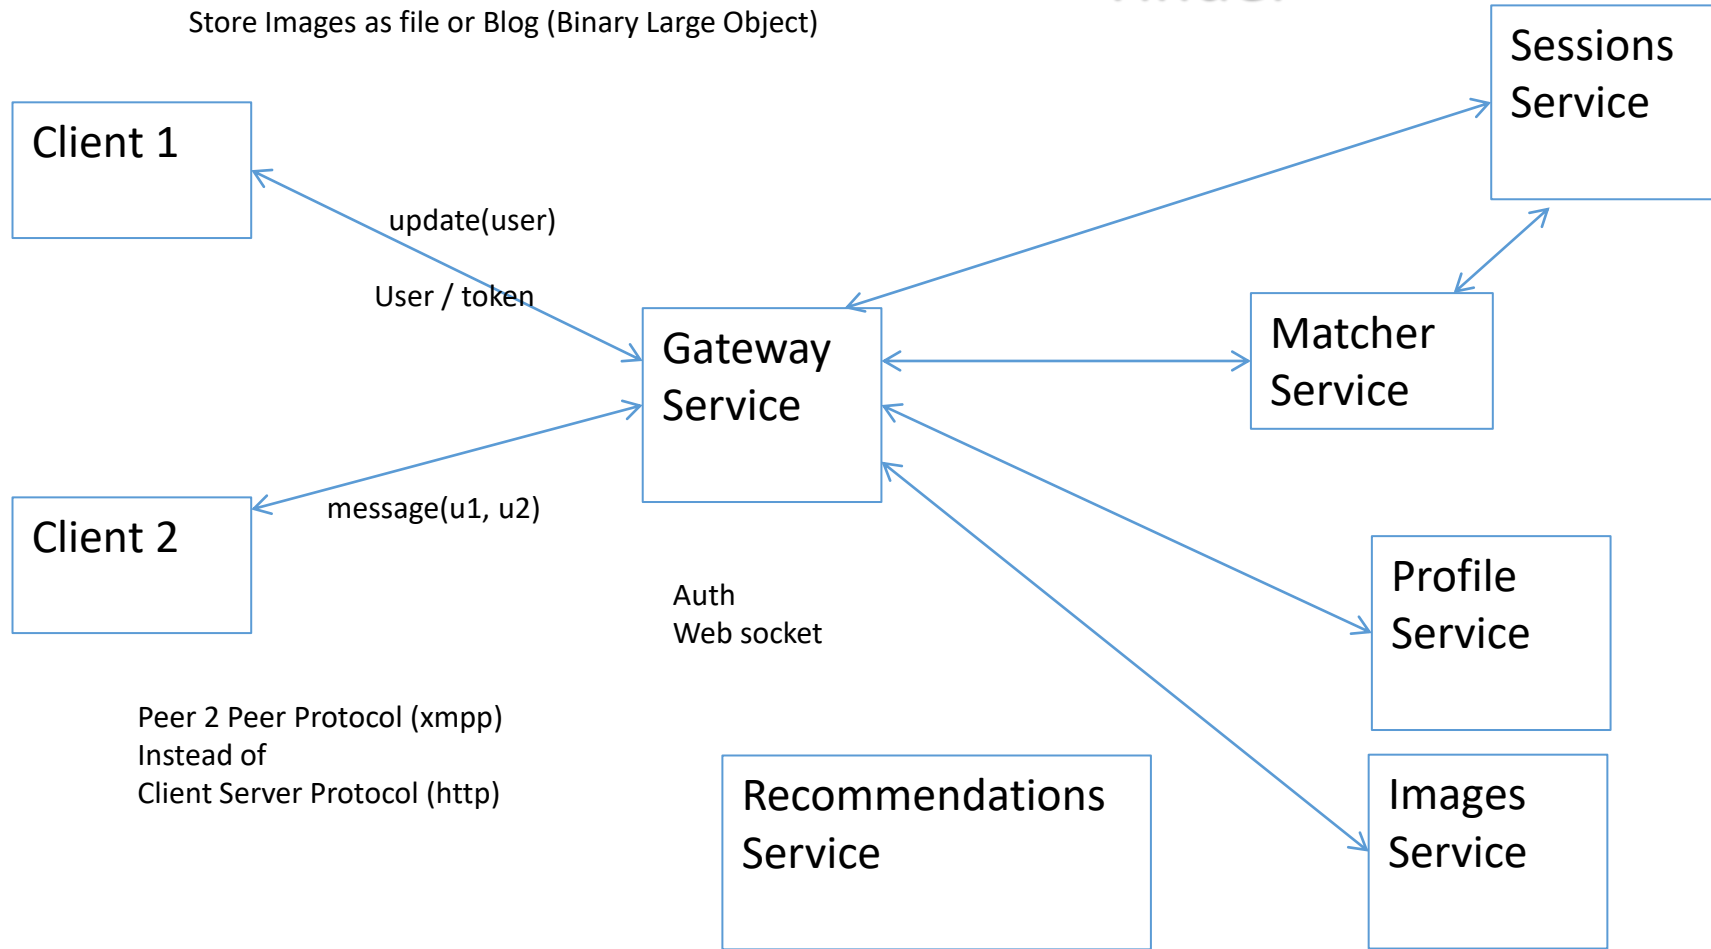
Database Sharding

- Consistency
- Availability
- Can shard by userId, location
- Problems
- Joins across shards, expensive
- Shards are inflexible in number, can use consistent hashing (Memcached)
- Create index on shards
- Master -> Slave Arch
- Writes to master, reads to slaves, slave becomes master if needed

Single point of failure

- More Nodes
- Master – Slave
- Regions
- Multiple Load Balancers (Gateway)
- Client -> DNS -> LBs -> Nodes -> DBs

Tinder



Can only sort by 1 Index

- 1) Age
- 2) Gender
- 3) Location

Use

- 1) Cassandra/ Dynamo
 - 2) Sharding, Horizontal Partitioning on location
- Use master slave arch

Features

- 1) Store Profiles (Images – 5 per User)
Store Images as file or Blob (Binary Large Object)
- 2) Recommend matches (# of active users)
- 3) Note matches (0.1% match)
- 4) Direct messaging

DBs

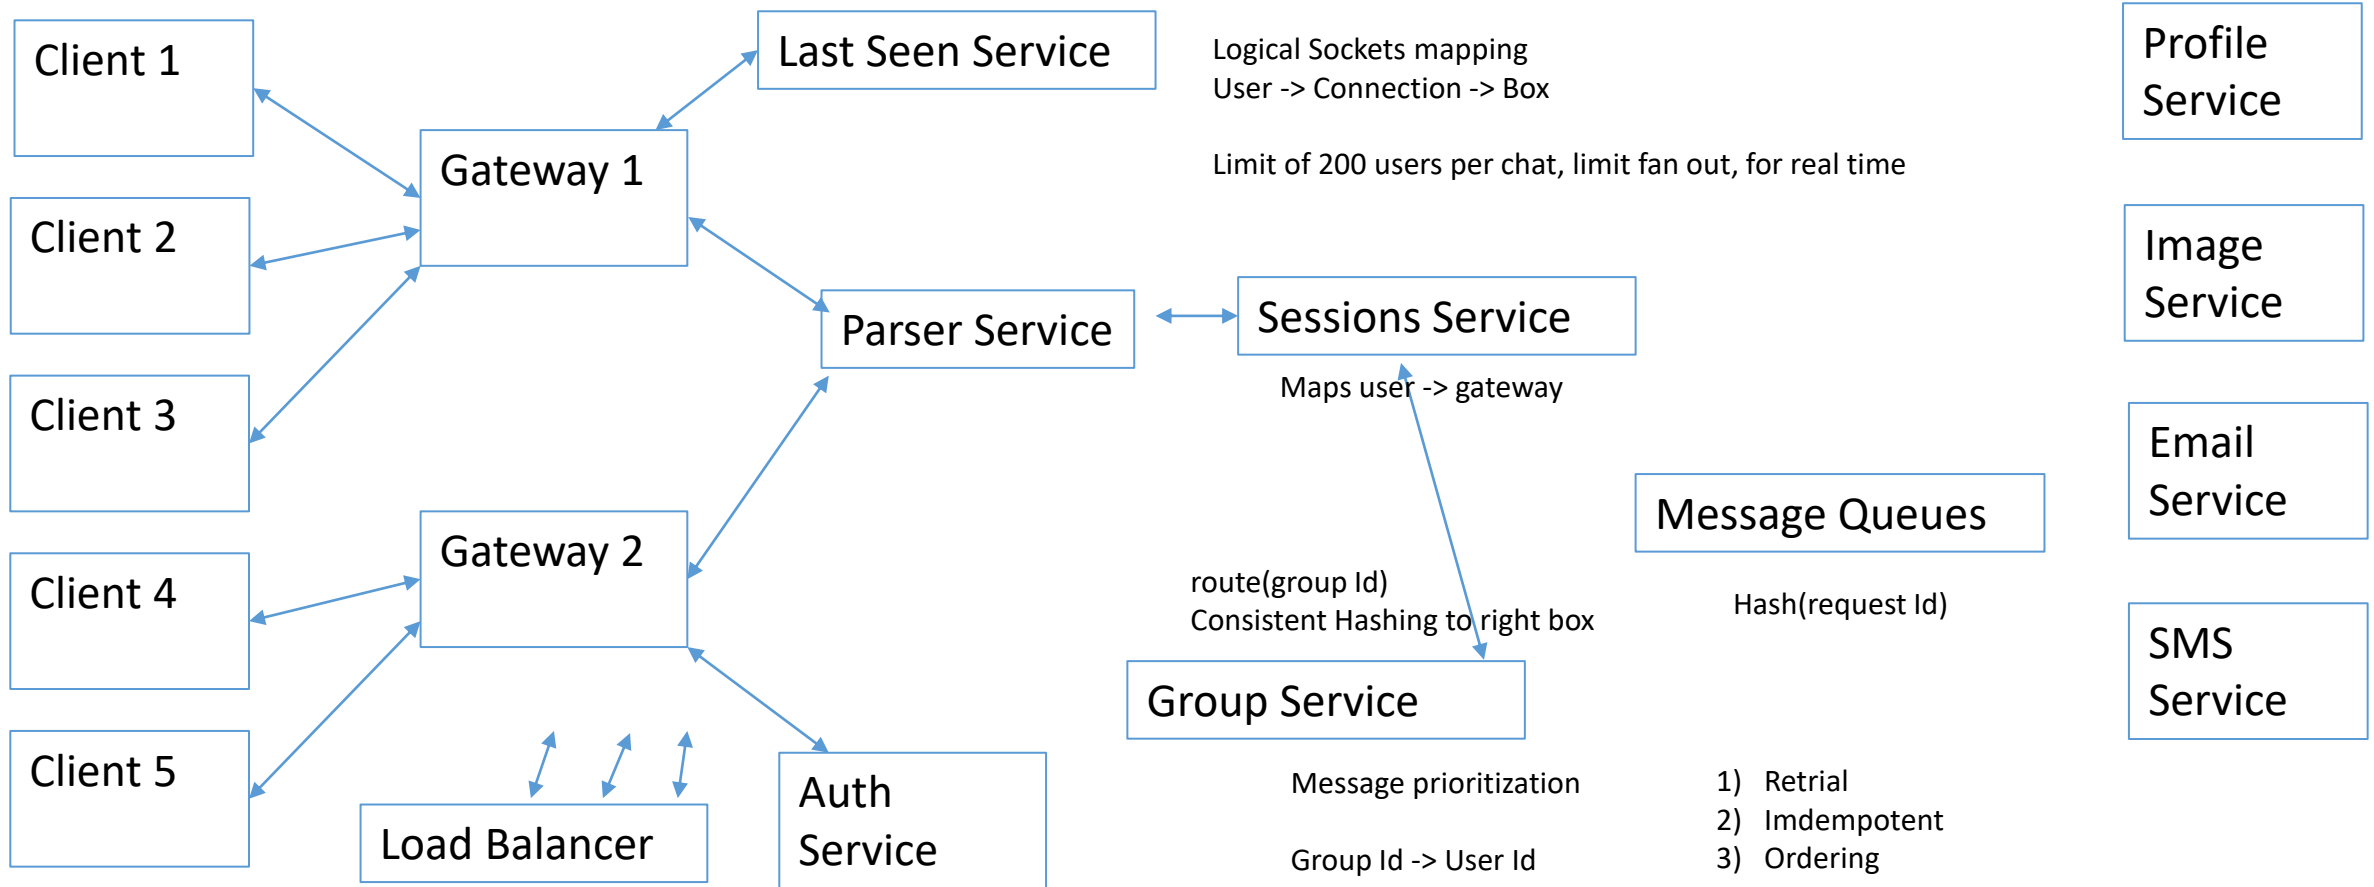
- 1) Mutability
- 2) Transaction – ACID
- 3) Indexes (search)
- 4) Access Control

Store as file

- 1) Cheaper
 - 2) Faster
- profileId, imageId, FileUrl
To Distributed File Sys

WhatsApp

Web sockets over http long polling



Features

- 1) Group Messaging
- 2) Sent, Delivered, Read Receipts
- 3) Online/ Last Seen
- 4) Image Sharing
- 5) Chats are temporary/ permanent

Deprioritize services like seen/read messages under huge loads

Distributed Caching

Why

- 1) Avoid Network Call
 - 2) Avoid computations
 - 3) Reduce DB load
- If cache closer to server, faster, but can be inconsistent between servers (in memory)
 - Global cache (Redis), can recover when failing, and can scale independently
 - Write-through
 - Update cache first, then update DB
 - Can wait and send to DB in both, for noncritical data, for saving network calls
 - Not practical for multiple caches
 - Write-back (performance issues)
 - Update DB first, then make entry in cache/ or invalidate
 - When hit cache on GET, then go to db, and also update cache
 - Hit cache, invalidate entry if it is there
 - Hybrid
 - If not critical info, write to cache, wait, and take entries in bulk and write to DB
- 1) Eviction Policy
 - 2) Thrashing – constantly inputting and outputting to cache without using results
 - 3) Consistency

API Design

Items of importance

- 1) Naming
- 2) Parameters deficiency
 - 1) More params for efficiency only
- 3) Response object simplicity
- 4) Return specific errors messages, for expected errors

Design API request

- 1) www.webiste.com/chat_messaging/getAdmins
 - 1) POST
 - 2) Request Object
 - 3) Response
- 2) No side effects. If many flags, should break down into separate functions
 - 1) Doing everything
 - 2) atomicity
- 3) When response is huge
 - 1) Pagination - Break response into multiple responses, but not stateless
 - 2) Fragmentation
- 4) Do you want perfect consistency?

NoSQL

1) SQL

- 1) requires joins which are expensive, no way to efficiently normalize

2) NoSQL


- 1) Flexible schema
- 2) insertions and retrievals require the whole blob
- 3) Horizontal partitioning built in, build for scale
- 4) Build for aggregation, finding metrics

3) NoSQL – disadvantages



- 1) Not built for updates (delete and insert), Consistency, ACID not guaranteed
- 2) Not read optimized
- 3) Relations are not implicit
- 4) Joins are hard

4) Cassandra

- 1) Have n instances
- 2) Writes to m servers, where pos through $pos + m - 1$
- 3) Quorum factor of x . For reads, x servers have to return the same value

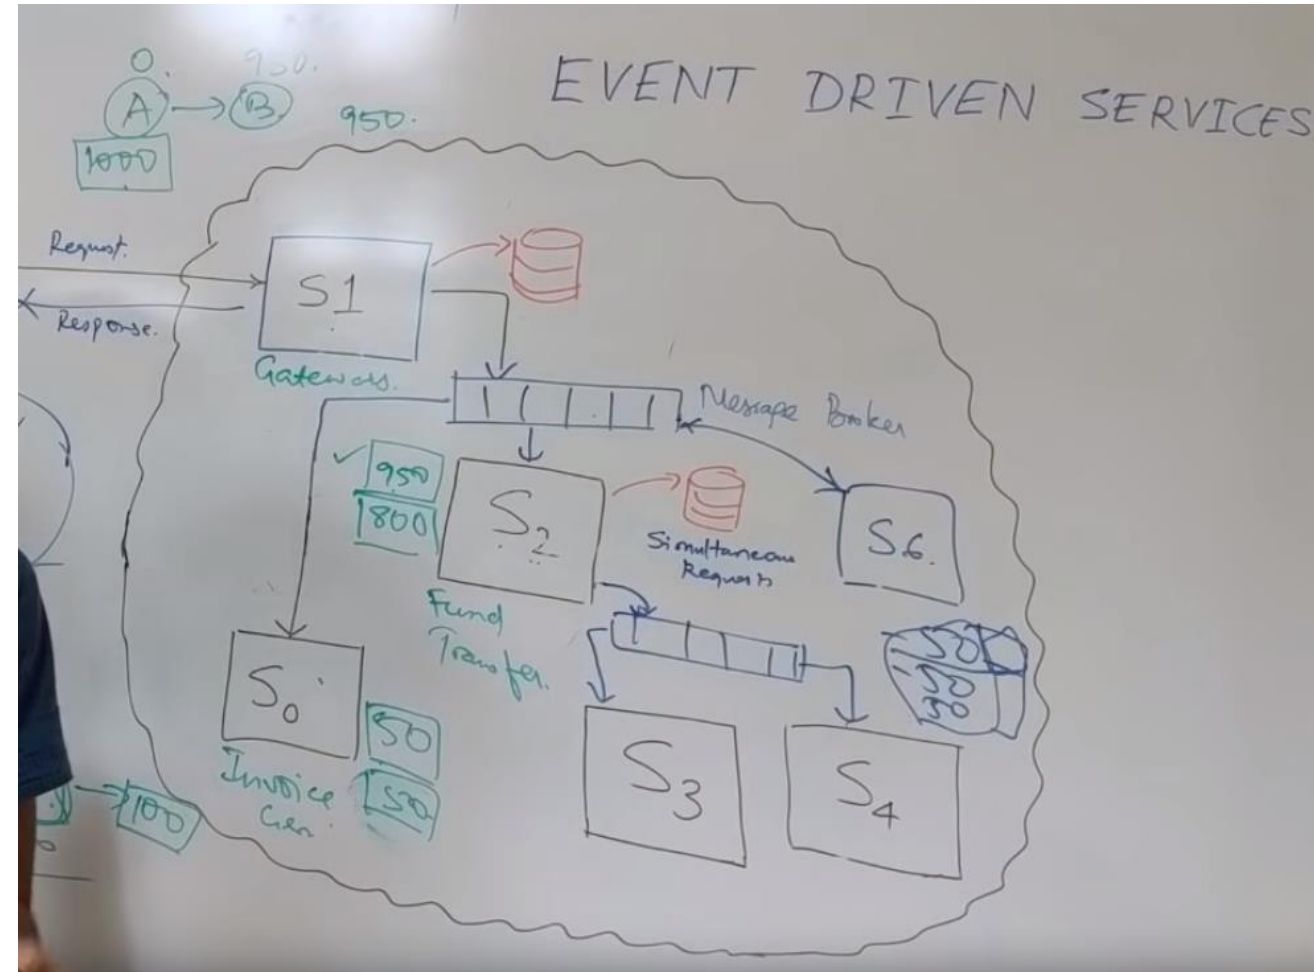
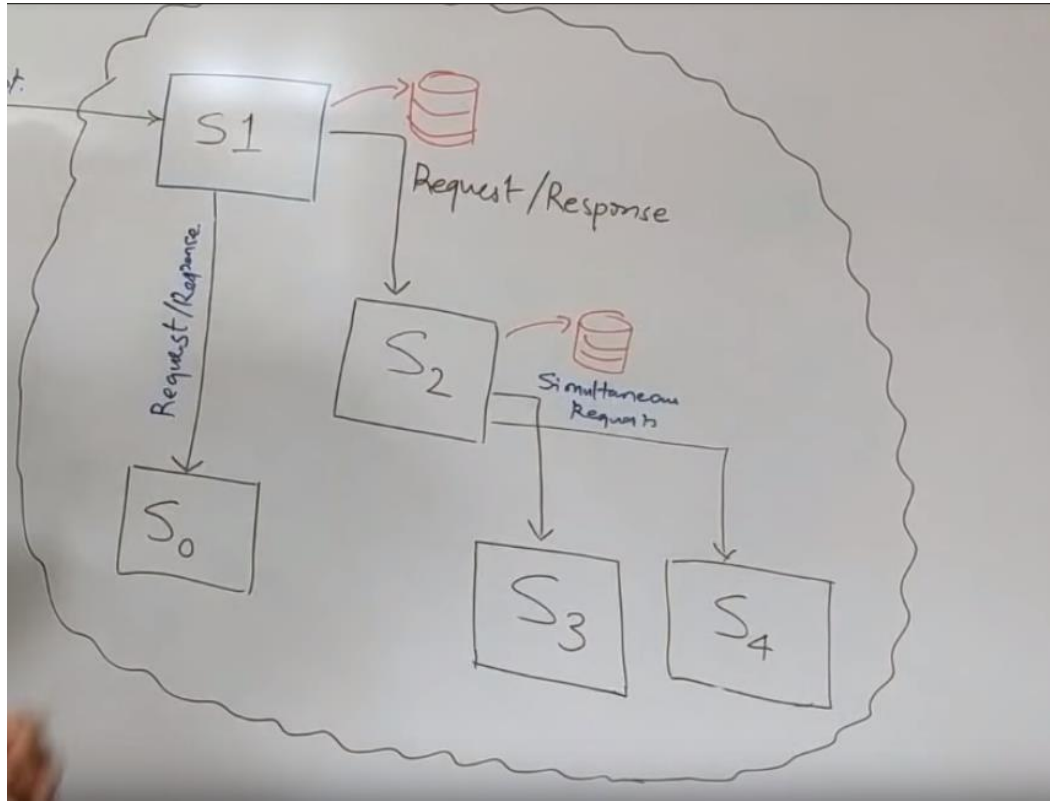


Distributed Consensus



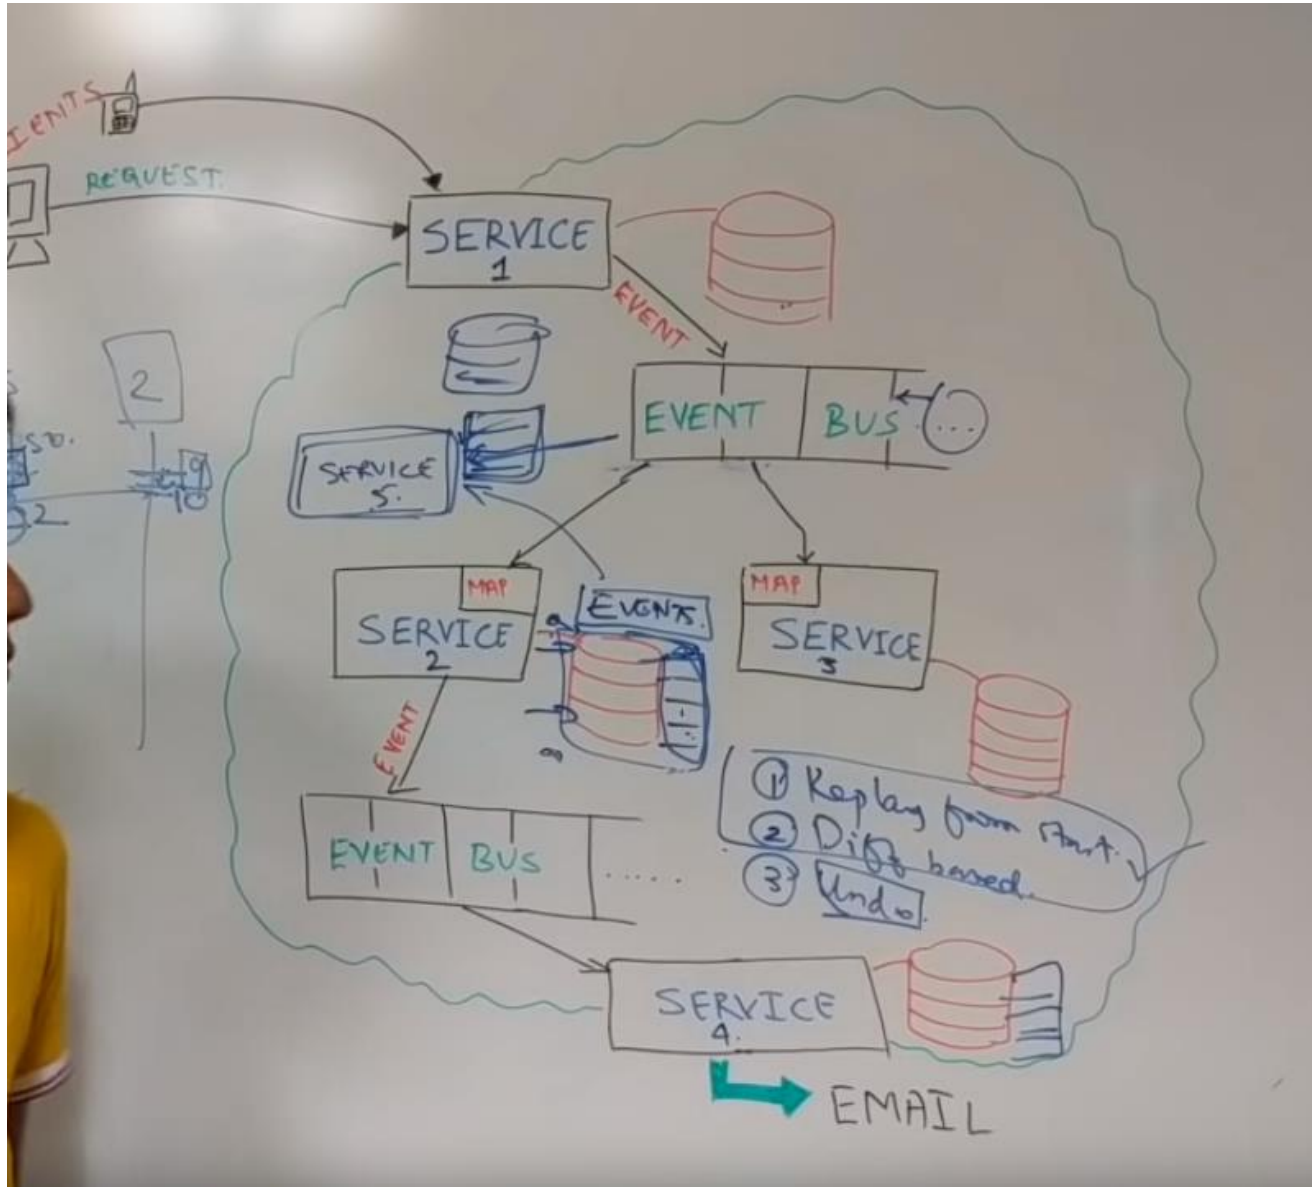
- 1) Bottleneck to Master DB
 - 1) Add slave node, prevent writes to it, only reads
 - 2) Or make both masters
- 2) Add 2 slaves
- 3) 2 Phase commit, 3 phase commit, MVCC, SAGA

Pub Sub



- 1) Use queue
- 2) Disadvantage – not enough consistency by default

Event Driven Systems



- 1) Availability (but lower consistency)
- 2) Easy Roll- backs
- 3) Replacement

Problems

- 1) Consistency
- 2) N/A to Gateways
- 3) Less control of responses
- 4) Compaction
- 5) Hidden Slow
- 6) Fixing
 - 1) Replay from Start
 - 2) Diff based
 - 3) Undo

Instagram

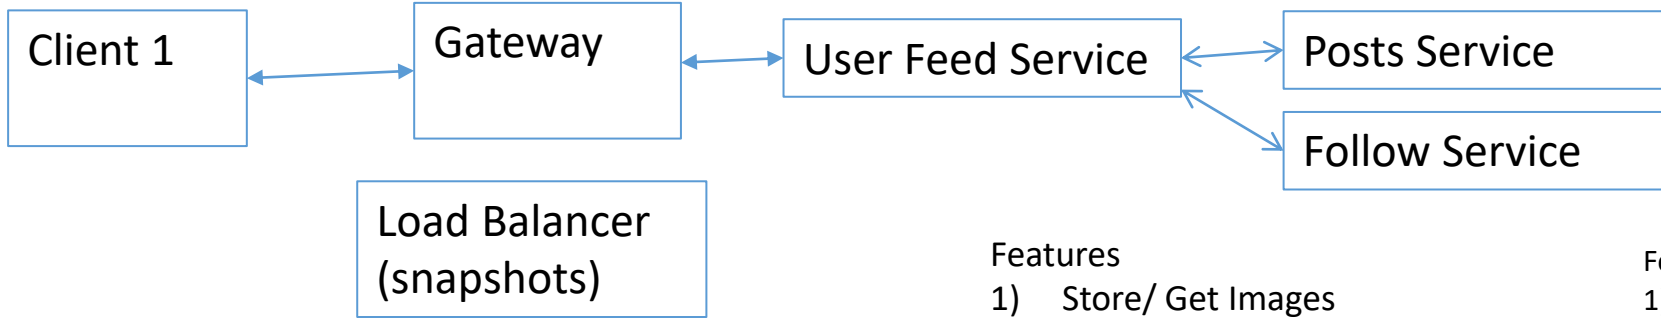


Image Service

Chat Service

Profile Service

Activity Service

Sessions Service

Feature #4
Publish Newsfeed

Features

- 1) Store/ Get Images
- 2) Like + Comment on Images
- 3) Follow someone
- 4) Publish a newsfeed

Feature #1

- 1) Store/ Get Images

Already cover by Tinder slide

Get followers, then posts

- 1) Expose API
 - 1) Get posts by users
- 2) Limit # of posts from post service
- 3) Pre-compute User feed
post service sends event
user feed service, get followers,
add to each one's queue
 - 1) Store in cache
- 4) Celebrity posts
 - 1) Push Batch processing,
send 1000
posts every 10 sec (push)
 - 2) Wait for followers to poll
server
- 5) Hybrid model – push for normal user,
poll for celebrity

Feature #2

Like + Comment on Images

Can you comment on a comment?
Can you like a comment?
Design ER diagram and DB tables

Post				
ID	text	imageUrl	userId	Time

Comment			
ID	text	time	activityid

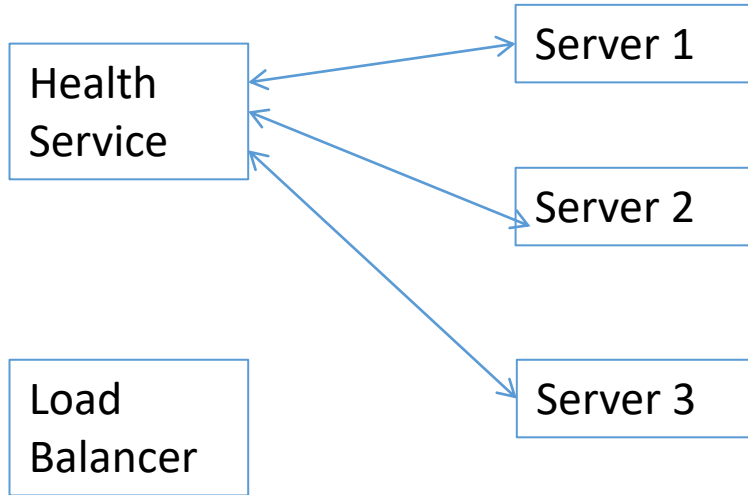
Likes					
ID	type	active	activityId	UserID	Time

Activity	
ID	numLikes

Feature #3
Follow someone

Follow		
followerID	FolloweeID	time

Heartbeats



If miss one heartbeat,

Mark as critical

If miss 2 heartbeats,

Ask another service to restart


1) Try 2 way heartbeat, to prevent zombie processes

1) Can restart itself

2) Cache snapshots on load balancer

1) Health services keep track of diffs

Tips



- 1) Don't go into detail prematurely
 - 1) Look for first point to go into detail
 - 2) Database -> ER diagrams
- 2) Do not have a set architecture in mind
- 3) KISS
 - 1) Keep It Simple Stupid
- 4) Form your thoughts
 - 1) Make points without justifications
 - 2) i.e. choosing a SQL database
- 5) Be Tech Aware